

# [ HFile Direct Read vs. HDFS SequenceFile ]

200,000 rows (value size: 500 – 8,000 bytes/row, 10 columns/row, no compression)

AWS EC2 m1.xlarge (4-core processor, 15 GB RAM) x 1 node

throughput  
[MB/sec]

**Note:** HDFS SequenceFile.Reader, without compression, doesn't do block-read but record-read, resulting relatively poor performance. Enabling block-read may help.

